

“Structured Query Language”

- It is NOT much like other programming languages
- It is NOT PROCEDURAL
- It does not process one record at a time, rather, it is a SET processing language
- All inputs to SQL are tables
- The output from a query is a table
- Output from a query referred to as the “Answer Set”
- Some queries may produce “interim” temporary answer sets

“Structured Query Language”

- It is a relatively simple language – brief syntax, few commands
- It is a relatively powerful language – a FEW lines of code can accomplish a LOT of work
- ANSI (American National Standards Institute) maintains a specification for “standard” SQL
- Each DBMS manufacturer follows the ANSI standard, but also adds extended features unique to their SQL

Useful Tool for Managing your Databases

- “SQLYog” from WebYog
- The community edition is free
- <https://github.com/webbyog/sql yog-community/wiki/Downloads>

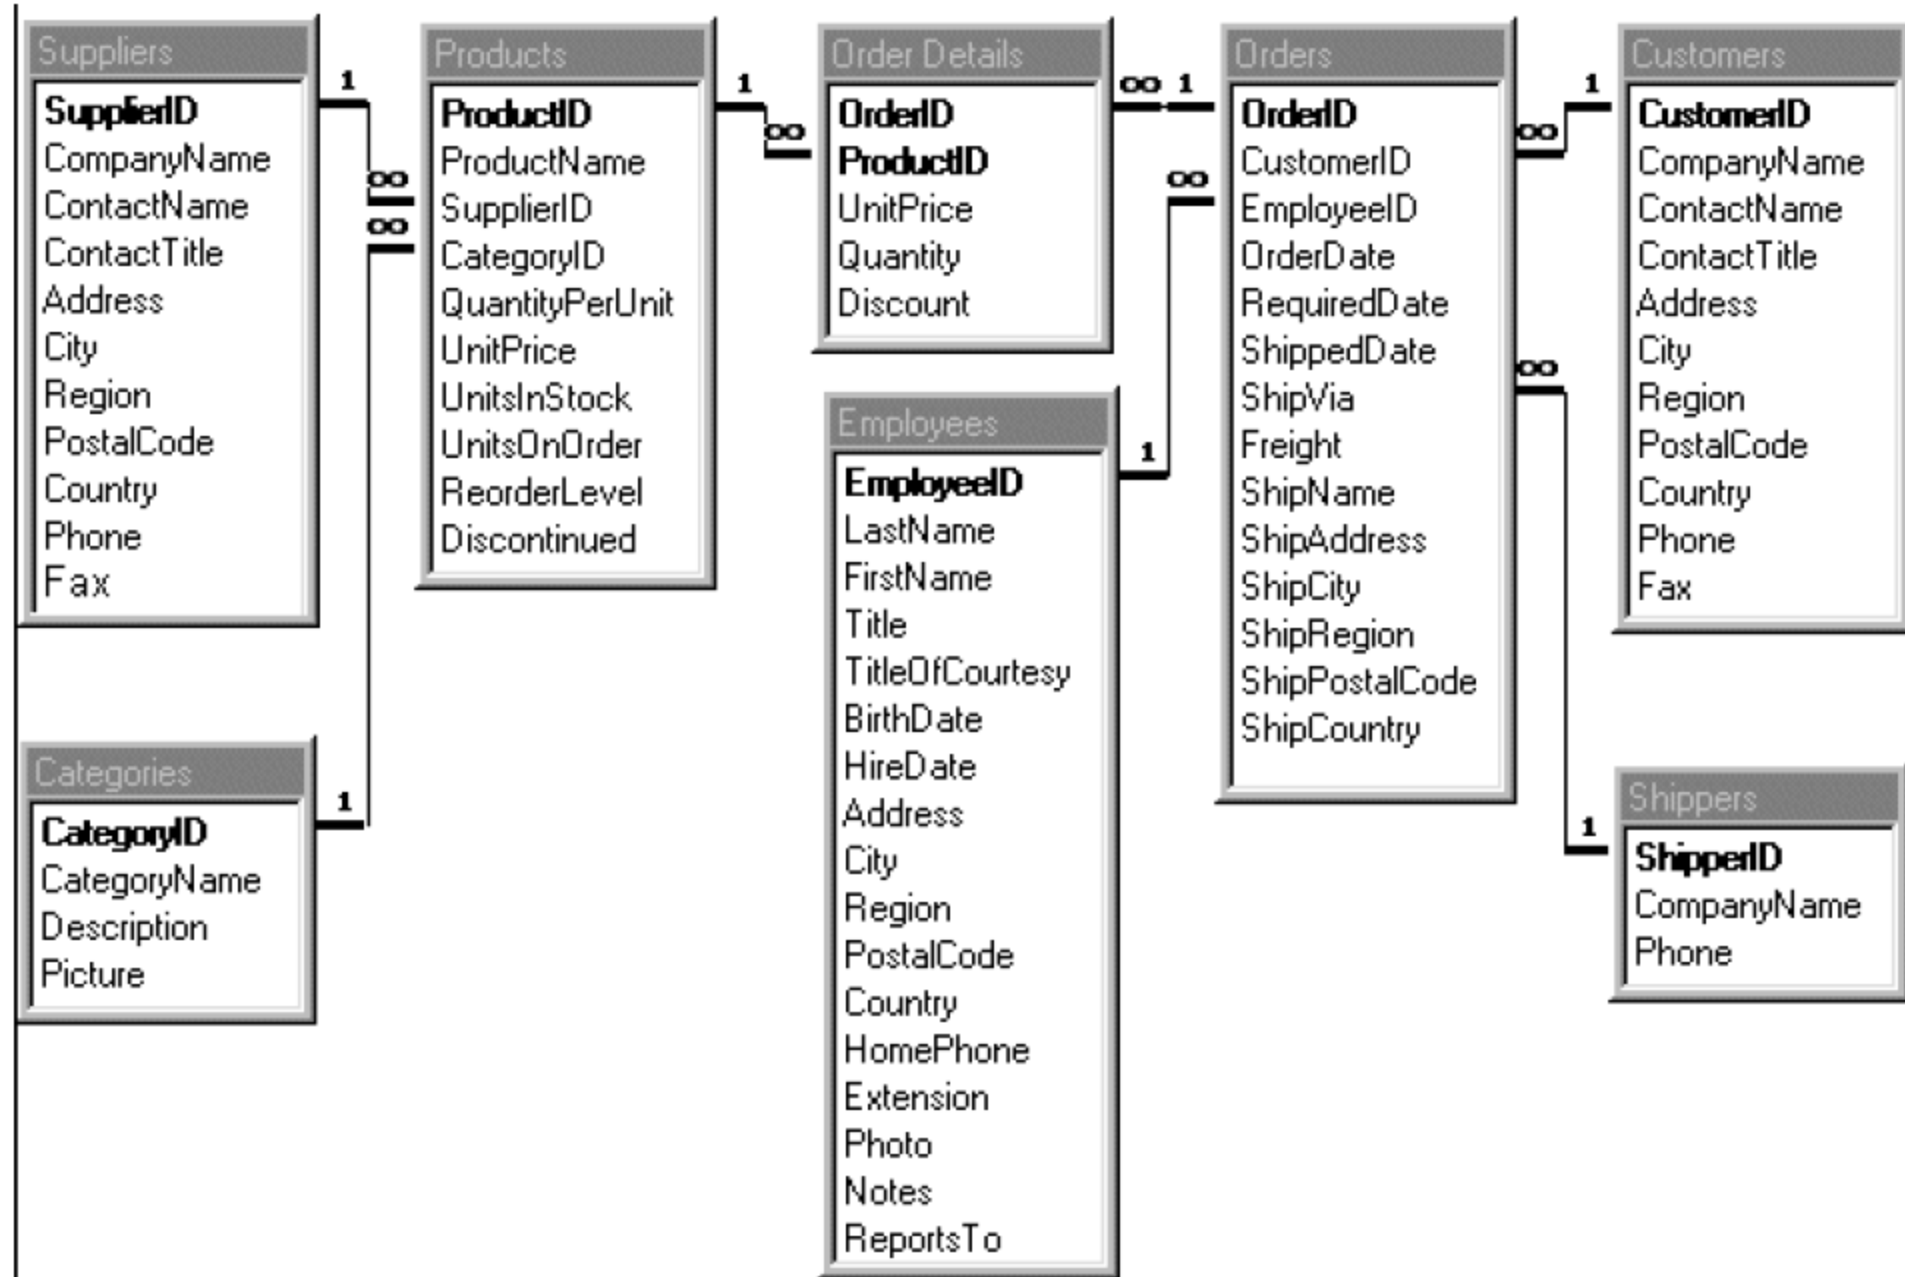
USE statement

```
USE <database>;
```

- Tells the Query Engine which database you want to use for your query

SELECT statement

```
SELECT <column1>, <column2>, <column3>,  
<literal>, <math expression>  
    FROM <table A> ;
```



Examples

```
select *  
    from nwEmployees;
```

```
use NorthWinds;
```

```
select EmployeeID, LastName, FirstName  
    from nwEmployees;
```

SELECT statement

- Literals may be either 'Character' (in quotes) or Numeric
- Math expressions

Only use with columns defined as numeric data types

+	Add
-	Subtract
*	Multiply
/	Divide
**	Exponent

SELECT statement

- Rename a column in the answer set with “AS”

```
Select employeeID as 'EMP'
```

- Concatenate character columns with (MySQL only)

```
Concat (<column1>, column2>)
```

- Comment out a line or part of a line of code by prefixing it with “- -” or embedding a “#”

- In Yog, highlight code then <ctrl>+<shift>+c to comment
<ctrl>+<shift>+r to remove comment

- Limit the size of the answer set with “limit” (MySQL)

```
Select LastName, FirstName  
from nwEmployees limit 100;
```

Examples

```
select 'Roster', LastName, FirstName  
      from nwEmployees;
```

```
select 'Roster' as 'Type', LastName, FirstName  
      from nwEmployees;
```

```
select 22, LastName, FirstName  
      from nwEmployees;
```

```
select 2 * 2, LastName, FirstName  
      from nwEmployees;
```

```
select concat(FirstName, ' ', LastName)  
      from nwEmployees;
```

SELECT statement

```
SELECT <column1>, <column2>, <column3>,  
<literal>, <math expression> AS <label>  
    FROM    <table A>  
    WHERE   <condition> ;
```

- The WHERE clause results in a subset of ROWs to appear in the answer set
- The condition in the WHERE clause takes this format:
 < operand > < operator > < operand >
- Operands may be columns or literals or expressions
- Operator may be
 - = Equals
 - <> Not equals
 - > Greater than
 - < Less than
 - Like
 - Between
 - In

- Operator may be: In or Like
In (literal, literal, literal)
Like 'string' with % or _ as a wildcard
- Multiple conditions may be joined with Boolean operators
AND, OR
- Conditions may be negated with Boolean operator
NOT
- Answer Set rows may be sorted with “Order By”
- Order By defaults to Ascending, can specify DESC

Distinct:

- The answer set may contain duplicate rows
- The “`distinct`” keyword before a column removes duplicates
- Example:
 - 87 Customers, each one has a country
 - How many distinct countries are they from?

Examples

```
select Customerid, ContactName, Region, Country
    from nwCustomers;
```

```
select Customerid, ContactName, Region, Country
    from nwCustomers
    where Country = 'Brazil';
```

```
select Customerid, ContactName, Region, Country
    from nwCustomers
    where Country <> 'Brazil';
```

```
select ProductID, ProductName, UnitPrice
    from nwProducts
    where UnitPrice > 60;
```

Examples

```
select ProductID, ProductName, UnitPrice
    from nwProducts
    where UnitPrice between 20 and 30;
```

```
select ProductID, ProductName, categoryid, UnitPrice
    from nwProducts
    where UnitPrice between 20 and 30
    and categoryid in (2, 4, 6);
```

```
select ProductID, ProductName, QuantityPerUnit
    from nwProducts
    where QuantityPerUnit like '%jars%';
```


Examples: Using distinct

```
Select CompanyName, ContactName, Country  
      from nwCustomers;
```

```
Select Country  
      from nwCustomers;
```

```
Select Distinct Country  
      from nwCustomers;
```

Handling Dates in MySQL

- MySQL supports DATE, DATETIME, and TIMESTAMP data types
- Columns with a data type of “TIMESTAMP” are stored as a 4-byte binary integer representing the number of seconds since 1970-01-01 00:00:00 UTC. TIMESTAMP has a range of '1970-01-01 00:00:01' UTC to '2038-01-19 03:14:07' UTC.
- If no value is provided for the TIME portion of a DATETIME column, it defaults to 00:00.00.0000
- To make it easier for humans to deal with date/time, MySQL allows us to reference dates/times in this format:
 YYYY-MM-DD and HH:MM.SS.nnn
- If you pass the date to MySQL as text in YYYY-MM-DD format, it will automatically convert it to the proper binary number
- If you pass the time to MySQL as text in HH:MM.SS.nnn format, it will automatically convert it to the proper binary number

Handling Dates in MySQL

YYYY-MM-DD and hh:mm:ss.nnn


- YYYY is four digits from 1000 through 9999 that represent a year.
- MM is two digits, ranging from 01 to 12, that represent a month in the specified year.
- DD is two digits, ranging from 01 to 31 depending on the month, that represent a day of the specified month.
- hh is two digits, ranging from 00 to 23, that represent the hour.
- mm is two digits, ranging from 00 to 59, that represent the minute.
- ss is two digits, ranging from 00 to 59, that represent the second.
- nnn is zero to three digits, ranging from 0 to 999, that represent the fractional seconds.

Examples: Using DATES

```
Select Now();  
Select Curdate();  
Select Curtime();  
Select Lastname, Firstname, Extract(Year From HireDate) AS HireYear  
    from NWEmployees;  
SELECT EmployeeID, Lastname, Firstname,  
    ROUND(DATEDIFF(HireDate, BirthDate)/365,0) AS HIRE_AGE  
FROM NWEmployees;
```

Examples: Using DATES

```
SELECT DATE_FORMAT(HireDate, '%b %d %Y %h:%i %p') FROM NWEMPLOYEES;
```

Format	Description
%a	Abbreviated weekday name (Sun-Sat)
%b	Abbreviated month name (Jan-Dec)
%c	Month, numeric (0-12)
%D	Day of month with English suffix (0th, 1st, 2nd, 3rd, )
%d	Day of month, numeric (00-31)
%e	Day of month, numeric (0-31)
%f	Microseconds (000000-999999)
%H	Hour (00-23)
%h	Hour (01-12)
%l	Hour (01-12)
%i	Minutes, numeric (00-59)
%j	Day of year (001-366)
%k	Hour (0-23)
%l	Hour (1-12)

SELECT statement

```
SELECT <column1>, <column2>, <column3>, <literal>,  
<math expression>  
    FROM <tableA>  
    WHERE <condition>  
ORDER BY <column1>, <column2> [DESC] ;
```

SQL provides the following GROUP FUNCTIONS

SUM – Provides the sum of the values in a column across many rows

AVG - Provides the average of the values in a column across many rows

COUNT – Provides a count of how many rows have a value in a column, counted across many rows

MIN - Provides the lowest value in a column across many rows

MAX - Provides the highest value in a column across many rows

GROUP FUNCTIONS

- SUM, AVG must only be used with NUMERIC columns
- MIN, MAX can be used with any data type
- COUNT can be used with any column, or with a (*) to simply count rows
- Group functions require SQL to create an interim answer set, and then process the group function against the interim answer set, delivering a final answer set that contains only the final total for the function. Always returns an integer value.
- When you combine a GROUP FUNCTION with a WHERE clause, keep in mind that the WHERE clause simply reduces the number of rows in the INTERIM answer set before the GROUP function does its calculation.

Examples –

```
select COUNT(*) as 'Total'
    from nwEmployees;
select COUNT(Distinct Country) as 'Countries'
    from nwCustomers;
select SUM(UnitPrice) as 'Total Price'
    from nwProducts;
select MAX(UnitPrice) as 'High Price'
    from nwProducts;
select MIN(UnitPrice) as 'Low Price'
    from nwProducts;
    where UnitPrice > 0
select AVG(UnitsInStock) as 'Average Inventory'
    from nwProducts;
```

GROUP BY

- Group functions process against an interim answer set to return a value across many rows.
- Using a GROUP BY clause enables SQL to provide subtotals. The GROUP BY tells SQL to perform the group function against a subset of rows in the interim answer set and provide a total for each subset of rows.

VERY IMPORTANT RULE

When using a GROUP BY every column in the SELECT statement must either be a GROUP FUNCTION or a COLUMN that you are grouping by.

Examples –

```
select Country, COUNT(*) as 'Total'
    from nwCustomers
    GROUP BY Country
select Country, COUNT(Country) as 'Total'
    from nwCustomers
    GROUP BY Country
```

Why are these the same?

```
select CustomerID, Country, COUNT(Country) as 'Total'
    from nwCustomers
    GROUP BY Country
```

Why is this incorrect?

More Examples –

From which supplier does Northwinds carry the most inventory?

```
select SupplierID, SUM(UnitsInStock) as 'Inventory'
  from nwProducts
 group by SupplierID
```

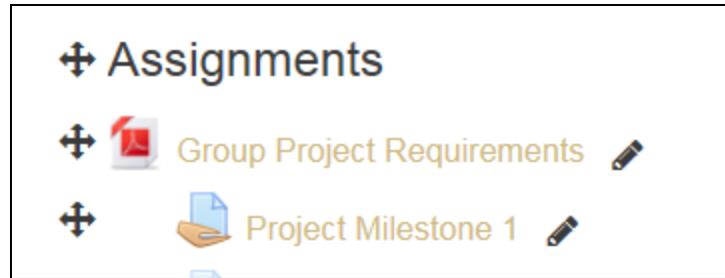
In which month in 2013 did Northwinds ship the most orders?

```
SELECT EXTRACT(MONTH FROM ShippedDate) AS 'Month',
       COUNT(OrderID) AS 'Orders'
  FROM nwOrders
 WHERE EXTRACT(YEAR FROM ShippedDate) = '2013'
 GROUP BY EXTRACT(MONTH FROM ShippedDate)
 ORDER BY 2 DESC
 Limit 1
```

Milestone 1:

- Link to submit is in Week Four (when the assignment was announced.)

- Looks like this:



- Due Date: Friday, October 6, 5:00 p.m. (One submission per team)
- Link to your github repo: email to your TA. (One submission per team)
- Yes, I changed the requirement last week and added a summary of your application architecture

HAVING

- Is simply like a WHERE clause against the answer set when you use a GROUP BY

Examples –

```
select Country, COUNT(*) as 'Total'
    from nwCustomers
    GROUP BY Country
```

In which countries does Northwinds have more than five customers?

```
select Country, COUNT(*) as 'Total'
    from nwCustomers
    GROUP BY Country
    HAVING COUNT(*) > 5
    order by 2 desc
```

SubQuery

Simply: a query within a query. The answer set to an “inner” query is used as a predicate in a where clause in the “outer” query.

- The subquery must return only one column.
- If the outer query WHERE clause contains an “equals” condition, the subquery must return ONE row.
- If the outer query WHERE clause contains an “in” condition, the subquery may return multiple rows, presented as a list of values.
- The Subquery is embedded within parentheses
- Outer and inner queries can hit two different tables

Examples –

```
select ProductID, ProductName, UnitPrice
  from nwProducts
 where UnitPrice = (

      select MAX(UnitPrice)
        from nwProducts )
```

Note that with the “equals” condition, the inner query returns only one value (one row, one column)

Examples –

```
select CustomerID, OrderID
  from nwOrders
 where OrderID in (
```

```
    select OrderID
      from nwOrderDetails
     where Quantity > 100 )
```

```
Order by CustomerID
```

- Note that with the “in” condition, the inner query returns many values (many rows, one column) as a list
- Uses Two different tables

Getting Data from Multiple Tables – The JOIN

- In order to run a query that retrieves data from multiple tables, those tables must be **JOINED**.
- Joining two tables requires that the two tables have a common key (typically a foreign key relationship) that appears in both tables.
- The common key columns need NOT have the same name, but must be of the same data type and length.
- A JOIN is one of the most **resource intensive** activities one can do in a relational database.

Basic Example

- Let's join nwOrders to nwEmployees
- nwOrders has 830 rows, each with an EmployeeID
- nwEmployees has 9 rows, each with an EmployeeID
- They have a common key: EmployeeID (primary key in nwEmployees; foreign key in nwOrders)
- We want SQL to join the rows in nwEmployees and nwOrders where the EmployeeID matches

Provide a listing showing Northwinds employees, sorted by LastName, and a count of each employee's orders

```
Select LastName, Firstname, count (OrderID) as 'Orders'
      from nwEmployees, nwOrders
     where nwEmployees.EmployeeID =
           nwOrders.EmployeeID
      GROUP BY LastName, FirstName
     Order By 1
```

Qualifying the Column Names

- Since the column “EmployeeID” exists in BOTH tables in this query, when referring to EmployeeID, we need to tell SQL which one.
- Therefore, we suffix the table name in front of the column name separated by a “ . ”
- Failure to fully qualify the column name will result in an “ambiguous column” error

Alternative

To save some typing, we can define an “alias” for each table. We can temporarily – only for the duration of this query -- rename the nwEmployees table “E”, and rename the nwOrders table “O”.

```
Select LastName, Firstname, count(OrderID) as 'Orders'
      from nwEmployees E, nwOrders O
     where E.EmployeeID = O.EmployeeID
    GROUP BY LastName, FirstName
   Order By 1
```

Alternative

- SQL allows another syntax option for doing the JOIN.
- These queries are equivalent:

```
Select LastName, Firstname, count(OrderID) as 'Orders'
    from nwEmployees E, nwOrders O
    where E.EmployeeID = O.EmployeeID
    GROUP BY LastName, FirstName
    Order By 1
```

```
Select LastName, Firstname, count(OrderID) as 'Orders'
    from nwEmployees E JOIN nwOrders O
    on E.EmployeeID = O.EmployeeID
    GROUP BY LastName, FirstName
    Order By 1
```


Beware the Cartesian Product

- Named after René DesCartes, philosopher & mathematician
- Famous for “Cogito ergo sum”
- Product = one table multiplied by another table
- The JOIN often creates a product, then selects rows from the product where the keys match
- For example, let’s join nwOrders to nwEmployees
- nwOrders has 14 columns, 830 rows
- nwEmployees has 17 columns, 9 rows
- The Cartesian product has 31 (14+17) columns, and 7470 (830 * 9) rows – most of which are meaningless



Cartesian Product

- SQL must go through the Cartesian Product (which is an INTERIM answer set) row-by-row, and select only those rows where the EmployeeID from nwEmployees is equal to the EmployeeID from nwOrders
- Therefore, we must include the WHERE clause that describes this condition
- Failure to fully qualify a JOIN operation with a join condition that matches all necessary keys will cause your answer set to include part or all of the Cartesian Product (which is mostly meaningless)
- The JOIN requires SQL to do a lot of work which consumes a lot of disk I/O and memory (= expensive)

Cartesian Product

Create a table JoinCustomer:

```
DROP TABLE IF EXISTS JoinCustomer;
CREATE TABLE JoinCustomer (
    CustomerID CHAR(5),
    CompanyName VARCHAR(430),
    ContactName VARCHAR(30),
    Country VARCHAR(15)
);
```

Load it with data from nwCustomers where Country = 'USA'

```
INSERT INTO JoinCustomer
    (CustomerID, CompanyName, Contactname, Country)
SELECT CustomerID, CompanyName, ContactName, Country
FROM nwCustomers
    where Country = 'USA';
```

Loads 13 rows of data.

Cartesian Product

Create a table JoinOrder:

```
DROP TABLE IF EXISTS JoinOrder;
CREATE TABLE JoinOrder (
    OrderID    INT(11),
    CustomerID CHAR(5),
    OrderDate  DATE,
    ShipCountry VARCHAR(15)
);
```

Load it with data from nwOrders where ShipCountry = 'USA'

```
INSERT INTO JoinOrder
    (OrderID, CustomerID, Orderdate, ShipCountry)
SELECT OrderID, CustomerID, Orderdate, ShipCountry
    FROM nwOrders WHERE ShipCountry = 'USA';
```

Loads 122 rows of data.

Cartesian Product from an unqualified join:

```
SELECT c.customerid, companyname, contactname, country,  
       OrderID, o.CustomerID, Orderdate, shipcountry  
FROM joincustomer c, joinorder o
```

customerid	companyname	contactname	country	OrderID	CustomerID	Orderdate	shipcountry
GREAL	Great Lakes Food Market	Howard Snyder	USA	10262	RATTC	2013-07-22	USA
HUNGC	Hungry Coyote Import Store	Yoshi Latimer	USA	10262	RATTC	2013-07-22	USA
LAZYK	Lazy K Kountry Store	John Steel	USA	10262	RATTC	2013-07-22	USA
LETSS	Lets Stop N Shop	Jaime Yorres	USA	10262	RATTC	2013-07-22	USA
LONEP	Lonesome Pine Restaurant	Fran Wilson	USA	10262	RATTC	2013-07-22	USA
OLDWO	Old World Delicatessen	Rene Phillips	USA	10262	RATTC	2013-07-22	USA
RATTC	Rattlesnake Canyon Grocery	Paula Wilson	USA	10262	RATTC	2013-07-22	USA
SAVEA	Save-a-lot Markets	Jose Pavarotti	USA	10262	RATTC	2013-07-22	USA
SPLIR	Split Rail Beer & Ale	Art Braunschweiger	USA	10262	RATTC	2013-07-22	USA
THEBI	The Big Cheese	Liz Nixon	USA	10262	RATTC	2013-07-22	USA
THECR	The Cracker Box	Liu Wong	USA	10262	RATTC	2013-07-22	USA
TRAIH	Trails Head Gourmet Provisioners	Helvetius Nagy	USA	10262	RATTC	2013-07-22	USA
WHITC	White Clover Markets	Karl Jablonski	USA	10262	RATTC	2013-07-22	USA
GREAL	Great Lakes Food Market	Howard Snyder	USA	10269	WHITC	2013-07-31	USA
HUNGC	Hungry Coyote Import Store	Yoshi Latimer	USA	10269	WHITC	2013-07-31	USA
LAZYK	Lazy K Kountry Store	John Steel	USA	10269	WHITC	2013-07-31	USA
LETSS	Lets Stop N Shop	Jaime Yorres	USA	10269	WHITC	2013-07-31	USA
LONEP	Lonesome Pine Restaurant	Fran Wilson	USA	10269	WHITC	2013-07-31	USA
OLDWO	Old World Delicatessen	Rene Phillips	USA	10269	WHITC	2013-07-31	USA
RATTC	Rattlesnake Canyon Grocery	Paula Wilson	USA	10269	WHITC	2013-07-31	USA
SAVEA	Save-a-lot Markets	Jose Pavarotti	USA	10269	WHITC	2013-07-31	USA
SPLIR	Split Rail Beer & Ale	Art Braunschweiger	USA	10269	WHITC	2013-07-31	USA

Cartesian Product

- The Cartesian Product contains 1586 rows. (122 * 13)
- Every row of JoinCustomer was multiplied by every row of JoinOrder.
- Only the rows where C.CustomerID = O.CustomerID are meaningful.

Joining three or more tables

- Every PAIR of tables being joined must have a common key
- Every PAIR of common keys must have a condition stated in a WHERE clause or in the “ON” clause of the JOIN
- Otherwise, your JOIN is not fully qualified and will result in a Cartesian Product (meaningless output)

Examples – Joining three tables

Create a report showing each employee and the total value of their orders sorted from highest value to lowest. (Order Value = UnitPrice * Quantity for each item on the order.)

```
Select LastName, Firstname,  
       sum(UnitPrice * Quantity) as 'OrderValue'  
from   nwEmployees E  
JOIN   nwOrders O on E.EmployeeID = O.EmployeeID  
JOIN   nwOrderDetails D on O.OrderID = D.OrderID  
GROUP BY LastName, FirstName  
Order By 3 desc
```


Examples – Joining three tables

Same Query, Different Syntax

```
Select LastName, Firstname,  
       sum(UnitPrice * Quantity) as 'OrderValue'  
from nwEmployees E, nwOrders O, nwOrderDetails D  
where E.EmployeeID = O.EmployeeID  
      and O.OrderID = D.OrderID  
GROUP BY LastName, FirstName  
Order By 3 desc
```

- **Join types**

Explicit inner join

```
SELECT * FROM employee  
INNER JOIN department ON employee.DepartmentID =  
    department.DepartmentID
```

Implicit inner join:

```
SELECT * FROM employee, department  
WHERE employee.DepartmentID =  
    department.DepartmentID
```

- Join types (cont.)

Outer Join

```
SELECT * FROM employee  
    LEFT OUTER JOIN department ON  
    employee.DepartmentID =  
    department.DepartmentID
```

Returns ALL rows from LEFT table and only matching rows from RIGHT table.

- **Analysis using Outer Joins**

- Are there some rows in `nworders` that do not have a valid foreign key reference to `nwemployees`?
- Are there some customers that have no orders?

```
SELECT COUNT(customerid) FROM nwcustomers
```

- there are 87 customers in `nwcustomers`

```
SELECT COUNT(distinct customerid) FROM nworders
```

- there are 89 distinct customers in `nworders`

- What's the difference?

- We want to find the orders in nworders whose customerID is NOT in nwcustomers

Method One: use a subquery

```
SELECT DISTINCT customerID
FROM nworders
WHERE customerID NOT IN (
    SELECT customerID FROM nwcustomers);
```

Method Two: use an outer join

```
SELECT DISTINCT O.customerID
FROM nworders O LEFT OUTER JOIN nwcustomers C
ON O.customerID = C.customerID
WHERE C.customerID IS NULL
```

This shows us FOUR customers who have orders in nworders that have no matching row in nwcustomers

- We want to find the customers in nwcustomers who have no orders in nworders

Method One: use a subquery

```
SELECT customerID
FROM nwcustomers
WHERE customerID NOT IN (
    SELECT DISTINCT customerID FROM nworders);
```

Method Two: use an outer join

```
SELECT DISTINCT C.customerID
FROM nworders O RIGHT OUTER JOIN nwcustomers C
ON O.customerID = C.customerID
WHERE O.customerID IS NULL
```

This shows us TWO customers who have no orders in nworders

DDL = Data Definition Language

Some SQL commands are used to DEFINE or MODIFY the structures in the database.

Create

Alter

Drop

DML = Data Manipulation Language

Some SQL commands are used to MODIFY the data in the database

Update

Insert

Delete

CREATE statement

```
CREATE TABLE <table name>
(column    DATATYPE (L) ,
column    DATATYPE (L)  NOT NULL,
column    DATATYPE (L)  CONSTRAINT <constr
name> TYPE,
column    DATATYPE (L) )
TABLESPACE <tablespace name>;
```


ALTER statement

```
ALTER TABLE <table name>
```

```
    ADD COLUMN column DATATYPE (L) ,  
    CONSTRAINT <constr name> TYPE
```

```
ALTER TABLE <table name>
```

```
    DROP CONSTRAINT <constr name>
```

DROP statement

```
DROP TABLE <table name>
```

UPDATE statement

```
UPDATE <table name>  
SET column = <value>  
WHERE <condition>
```

NOTE: if the WHERE is omitted, ALL rows are updated.

DELETE statement

```
DELETE FROM <table name>  
WHERE <condition>
```

NOTE: if the WHERE is omitted, ALL rows are deleted.

INSERT statement

```
INSERT INTO <table name>
```

```
VALUES (value, value, value, value)
```

(must have a value or NULL for every column in the table)

```
INSERT INTO <table name> (column, column,  
    column)
```

```
VALUES (value, value, value)
```

(if no column/value is specified, NULL or default will be assigned)

Getting more help with SQL

Best Online Tutorial: <http://www.w3schools.com/sql/>